

HOMEWORK 3

CSC311 Fall 2019

Nicole Xin Yue Wang, #1004235339

1. Fitting a Naïve Bayes Model

a) Let N be the number of inputs

so then dataset $X = [x^{(1)}, x^{(2)}, \dots, x^{(N)}]^T$

So we have:

$$\begin{aligned} \ell(\theta, \pi) &= \sum_{i=1}^N \log p(x^{(i)}, t^{(i)} | \theta, \pi) = \sum_{i=1}^N \log \left\{ p(t^{(i)} | \pi) \prod_{j=1}^{784} p(x_j^{(i)} | t^{(i)}, \theta_j) \right\} \\ &= \sum_{i=1}^N \left[\underbrace{\log p(t^{(i)} | \pi)}_{\ell(\pi)} + \underbrace{\sum_{j=1}^{784} \log p(x_j^{(i)} | t^{(i)}, \theta_j)}_{\ell(\theta)} \right] \end{aligned}$$

$$\ell(\theta_j) = \sum_{i=1}^N \log p(x_j^{(i)} | t^{(i)}, \theta_j) = \sum_{i=1}^N \sum_{c=0}^9 t_c^{(i)} \left[x_j^{(i)} \log \theta_{jc} + (1 - x_j^{(i)}) \log (1 - \theta_{jc}) \right]$$

for pixel j and class c :

$$\begin{aligned} \frac{\partial \ell(\theta_j)}{\partial \theta_{jc}} &= \sum_{i=1}^N t_c^{(i)} \left(\frac{x_j^{(i)}}{\theta_{jc}} - \frac{1 - x_j^{(i)}}{1 - \theta_{jc}} \right) \\ &= \sum_{i=1}^N \frac{t_c^{(i)} (x_j^{(i)} - \theta_{jc})}{\theta_{jc} (1 - \theta_{jc})} \end{aligned}$$

$$\begin{aligned} \text{Set to zero: } \sum_{i=1}^N t_c^{(i)} x_j^{(i)} &= \sum_{i=1}^N t_c^{(i)} \theta_{jc} \\ \hat{\theta}_{jc} &= \frac{\sum_{i=1}^N t_c^{(i)} x_j^{(i)}}{\sum_{i=1}^N t_c^{(i)}} \end{aligned}$$

\Rightarrow counts the number of images with a white pixel on j in all c labeled images.

$$\begin{aligned} \ell(\pi) &= \sum_{i=1}^N \log p(t^{(i)} | \pi) \\ &= \sum_{i=1}^N \log \prod_{c=0}^9 \pi_c^{t_c^{(i)}} \\ &= \sum_{i=1}^N \sum_{c=0}^9 t_c^{(i)} \log \pi_c \\ &= \sum_{i=1}^N \left[\sum_{c=0}^8 t_c^{(i)} \log \pi_c + t_9^{(i)} \log (1 - \sum_{c=0}^8 \pi_c) \right] \end{aligned}$$

Let constraint be $\sum_{c=0}^9 \pi_c = 1$, F be our Lagrange Multiplier

$$\Rightarrow F = \sum_{i=1}^N \left[\sum_{c=0}^8 t_c^{(i)} \log \pi_c + t_9^{(i)} \log (1 - \sum_{c=0}^8 \pi_c) \right] - \lambda \left(\sum_{c=0}^9 \pi_c - 1 \right)$$

$$\text{for } j=0, \dots, 8: F_{\pi_j} = \sum_{i=1}^N \left(\frac{t_j^{(i)}}{\pi_j} - \frac{t_9^{(i)}}{\pi_9} \right) - \lambda$$

$$F_{\pi_9} = \sum_{i=1}^N \frac{t_9^{(i)}}{\pi_9} - \lambda$$

$$F_{\lambda} = \sum_{c=0}^9 \pi_c - 1$$

$$\begin{aligned} \text{set } F_{\pi_9} = 0 &\Rightarrow \frac{1}{\pi_9} \sum_{i=1}^N t_9^{(i)} = \lambda \\ \hat{\pi}_9 &= \frac{\sum_{i=1}^N t_9^{(i)}}{\lambda} \quad (1) \end{aligned}$$

$$\text{set } F_{\pi_j} = 0 \Rightarrow \frac{1}{\pi_j} \sum_{i=1}^N t_j^{(i)} - \frac{1}{\pi_9} \sum_{i=1}^N t_9^{(i)} = \lambda$$

$$\begin{aligned} \Rightarrow \frac{1}{\pi_j} \sum_{i=1}^N t_j^{(i)} - \frac{1}{\pi_9} \sum_{i=1}^N t_9^{(i)} &= \frac{1}{\pi_9} \sum_{i=1}^N t_9^{(i)} \\ \frac{1}{\pi_j} \sum_{i=1}^N t_j^{(i)} &= \frac{2}{\pi_9} \sum_{i=1}^N t_9^{(i)} \\ \frac{\pi_j}{\pi_9} &= \frac{\sum_{i=1}^N t_j^{(i)}}{2 \sum_{i=1}^N t_9^{(i)}} \quad (2) \end{aligned}$$

$$\begin{aligned} F_{\lambda} &= -\pi_9 \sum_{c=0}^8 \frac{\pi_c}{\pi_9} + 1 \\ &= -\frac{\sum_{i=1}^N t_9^{(i)}}{\lambda} \left[\sum_{c=0}^8 \frac{\sum_{i=1}^N t_c^{(i)}}{2 \sum_{i=1}^N t_9^{(i)}} + 1 \right] + 1 \\ &= -\frac{1}{2\lambda} \sum_{i=1}^N \sum_{c=0}^8 t_c^{(i)} - \frac{\sum_{i=1}^N t_9^{(i)}}{\lambda} + 1 \end{aligned}$$

#1. a) continue...

$$\text{set } F_\lambda = 0$$

$$\Rightarrow 0 = \sum_{i=1}^N \sum_{c=0}^9 t_c^{(i)} + 2 \sum_{i=1}^N t_9^{(i)} - 2\lambda$$

$$0 = N + \sum_{i=1}^N t_9^{(i)} - 2\lambda \quad \# \sum_{c=0}^9 t_c^{(i)} = 1$$

$$\lambda = \frac{1}{2} (N + \sum_{i=1}^N t_9^{(i)})$$

$$\text{plug } \lambda \text{ into } \textcircled{1}: \quad \hat{\pi}_9 = \frac{2 \sum_{i=1}^N t_9^{(i)}}{N + \sum_{i=1}^N t_9^{(i)}}$$

$$\text{plug } \hat{\pi}_9 \text{ into } \textcircled{2}: \quad \hat{\pi}_j = \frac{2 \sum_{i=1}^N t_j^{(i)}}{N + \sum_{i=1}^N t_9^{(i)}} \cdot \frac{\sum_{i=1}^N t_j^{(i)}}{2 \sum_{i=1}^N t_9^{(i)}}$$

$$= \frac{\sum_{i=1}^N t_j^{(i)}}{N + \sum_{i=1}^N t_9^{(i)}}$$

b) By Bayes' Rule:

$$P(t^{(i)} | x^{(i)}, \theta, \pi) = \frac{P(x^{(i)} | t^{(i)}, \theta, \pi)}{P(x^{(i)} | \pi)} \frac{P(t^{(i)} | \pi)}{\prod_{j=1}^{784} P(x_j^{(i)} | t^{(i)}, \theta_{j,t^{(i)}})}$$

$$= \frac{\sum_{c=0}^9 P(x^{(i)} | c, \theta, \pi) P(c | \theta, \pi)}{\pi_{t^{(i)}} \prod_{j=1}^{784} P(x_j^{(i)} | t^{(i)}, \theta_{j,t^{(i)}})}$$

$$= \frac{\sum_{c=0}^9 \pi_c \prod_{j=1}^{784} P(x_j^{(i)} | c, \theta_{jc})}{\pi_{t^{(i)}} \prod_{j=1}^{784} \theta_{j,t^{(i)}}^{\alpha_j^{(i)}} (1 - \theta_{j,t^{(i)}})^{1 - \alpha_j^{(i)}}}$$

$$= \frac{\pi_{t^{(i)}} \prod_{j=1}^{784} \theta_{j,t^{(i)}}^{\alpha_j^{(i)}} (1 - \theta_{j,t^{(i)}})^{1 - \alpha_j^{(i)}}}{\sum_{c=0}^9 [\pi_c \prod_{j=1}^{784} \theta_{jc}^{\alpha_j^{(i)}} (1 - \theta_{jc})^{1 - \alpha_j^{(i)}}]}$$

$$\therefore \log P(t^{(i)} | x^{(i)}, \theta, \pi) = \log \pi_{t^{(i)}} + \sum_{j=1}^{784} [\alpha_j^{(i)} \log \theta_{j,t^{(i)}} + (1 - \alpha_j^{(i)}) \log (1 - \theta_{j,t^{(i)}})]$$

$$- \log \sum_{c=0}^9 [\pi_c \prod_{j=1}^{784} \theta_{jc}^{\alpha_j^{(i)}} (1 - \theta_{jc})^{1 - \alpha_j^{(i)}}]$$

Code for log_likelihood function:

```
def log_likelihood(images, theta, pi):
    image_num = np.shape(images)[0]
    class_num = np.shape(theta)[1]

    log_like = np.zeros((image_num, class_num))
    for i in range(image_num):
        for c in range(class_num):
            nume = pi[c] * np.prod(np.power(theta[:,c], images[i]) *
                                   np.power(1 - theta[:,c], 1 - images[i]))
            if (nume != 0):
                denom = 0
                for cc in range(class_num):
                    denom += pi[cc] * np.prod(np.power(theta[:,cc], images[i]) *
                                                np.power(1 - theta[:,cc], 1 - images[i]))
                log_like[i][c] = np.log(nume / denom)
            else:
                log_like[i][c] = 0

    return log_like
```

c) fitting θ and π using MLE

```
def train_mle_estimator(train_images, train_labels):
    (image_num, pixel_num) = np.shape(train_images)
    class_num = np.shape(train_labels)[1]

    theta_mle = np.zeros((pixel_num, class_num))
    pi_mle = np.zeros(class_num)


    for c in range(class_num):
        tc = train_labels[:, c].copy()
        for j in range(pixel_num):
            xj = train_images[:, j].copy()
            theta_mle[j,c] = np.sum(tc * xj) / np.sum(tc)
        if(c == 9):
            pi_mle[c] = 2 * np.sum(tc) / (image_num + np.sum(train_labels[:, 9]))
        else:
            pi_mle[c] = np.sum(tc) / (image_num + np.sum(train_labels[:, 9]))

    return theta_mle, pi_mle
```

Problem: there exist cases where θ_{jc} is 0, then it will be unable to find the log-likelihood since log of zero is N/A. If I set log-likelihood for where θ_{jc} is 0, I will get:

Average log-likelihood for MLE is -3.314481563095963

d) 10 grayscale images of MLE estimator plotted, one for each class



e) $\log p(\theta, x | \pi) = \log [p(\theta) p(x^{(i)} | \theta, \pi)] = \log p(\theta) + l_{MLE}$

Using Beta(3,3) as prior:

$$p(\theta, 3, 3) \propto \theta^2 (1-\theta)^2$$

$$\Rightarrow \log p(\theta) = 2 \log \theta + 2 \log(1-\theta)$$

$$\log p(\theta, x | \pi) = 2 \log \theta + 2 \log(1-\theta) + l_{MLE}$$

$$\begin{aligned} \frac{\partial}{\partial \theta_{jc}} \log p(\theta, x | \pi) &= \frac{2}{\theta_{jc}} - \frac{2}{1-\theta_{jc}} + \sum_{i=1}^N \frac{t_c^{(i)} (x_j^{(i)} - \theta_{jc})}{\theta_{jc} (1-\theta_{jc})} \\ &= \frac{2 - 4\theta_{jc}}{\theta_{jc} (1-\theta_{jc})} + \sum_{i=1}^N \frac{t_c^{(i)} (x_j^{(i)} - \theta_{jc})}{\theta_{jc} (1-\theta_{jc})} \\ &= \frac{2 - 4\theta_{jc} + \sum_{i=1}^N t_c^{(i)} (x_j^{(i)} - \theta_{jc})}{\theta_{jc} (1-\theta_{jc})} \end{aligned}$$

$$\hookrightarrow \text{set to zero} \Rightarrow 2 - 4\theta_{jc} + \sum_{i=1}^N t_c^{(i)} x_j^{(i)} - \theta_{jc} \sum_{i=1}^N t_c^{(i)} = 0$$

$$\hat{\theta}_{jc} = \frac{2 + \sum_{i=1}^N t_c^{(i)} x_j^{(i)}}{4 + \sum_{i=1}^N t_c^{(i)}}$$

f) fitting θ and π using MAP and finding accuracy.

```
def train_map_estimator(train_images, train_labels):
    (image_num, pixel_num) = np.shape(train_images)
    class_num = np.shape(train_labels)[1]
    theta_map = np.zeros((pixel_num, class_num))
    pi_map = np.ones(class_num)

    for c in range(class_num):
        tc = train_labels[:, c].copy()
        for j in range(pixel_num):
            xj = train_images[:, j].copy()
            theta_jc = (2 + np.sum(tc * xj)) / (4 + np.sum(tc))
            theta_map[j][c] = theta_jc
        if(c == 9):
            pi_map[c] = 2 * np.sum(tc) / (image_num + np.sum(train_labels[:, 9]))
        else:
            pi_map[c] = np.sum(tc) / (image_num + np.sum(train_labels[:, 9]))
    return theta_map, pi_map

def predict(log_like):
    image_num = np.shape(log_like)[0]
    predictions = np.zeros(np.shape(log_like))
    for i in range(image_num):
        max_index = np.argmax(log_like[i])
        predictions[i][max_index] = 1
    return predictions


def accuracy(log_like, labels):
    predictions = predict(log_like)
    match_num = 0
    N = np.shape(labels)[0]
    for i in range(N):
        if(np.array_equal(predictions[i], labels[i])):
            match_num += 1
    accuracy = match_num / N
    return accuracy
```

Average log-likelihood for MAP is -3.3669909342818136

Training accuracy for MAP is 0.83085

Test accuracy for MAP is 0.809

g) 10 greyscale images of MAP estimator plotted, one for each class



2. Generating from a Naïve Bayes Model

- a) True
- b) False
- c) Producing random images

```
def image_sampler(theta, pi, num_images):  
    (pixel_num, class_num) = np.shape(theta)  
    sampled_images = np.zeros((num_images, pixel_num))  
    cs = np.random.choice(class_num, num_images, p=pi)  
    for i in range(num_images):  
        sampled_images[i] = np.random.binomial(1, p=theta[:, cs[i]])  
    print("The random classes are: " + str(cs))  
    return sampled_images
```

Output:

The random classes are: [7 0 7 9 8 3 6 8 2 8]

